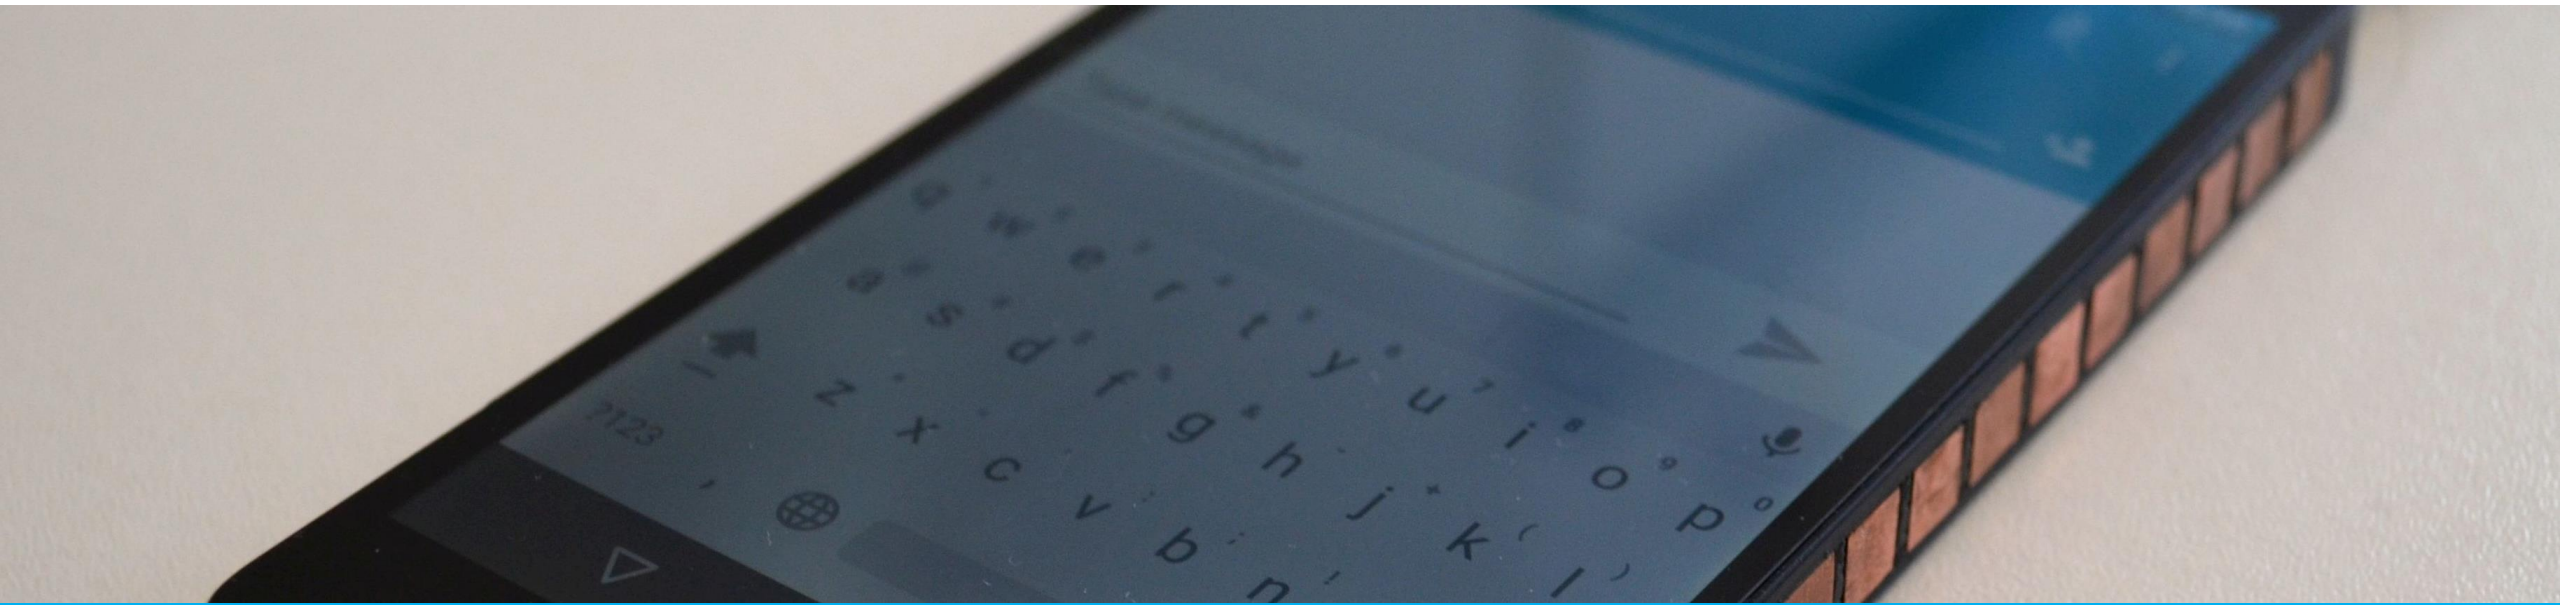




**University of Stuttgart**  
Institute for Visualization and Interactive Systems



# TensorFlow Mobile: Exporting and Optimizing Models for Mobile Devices

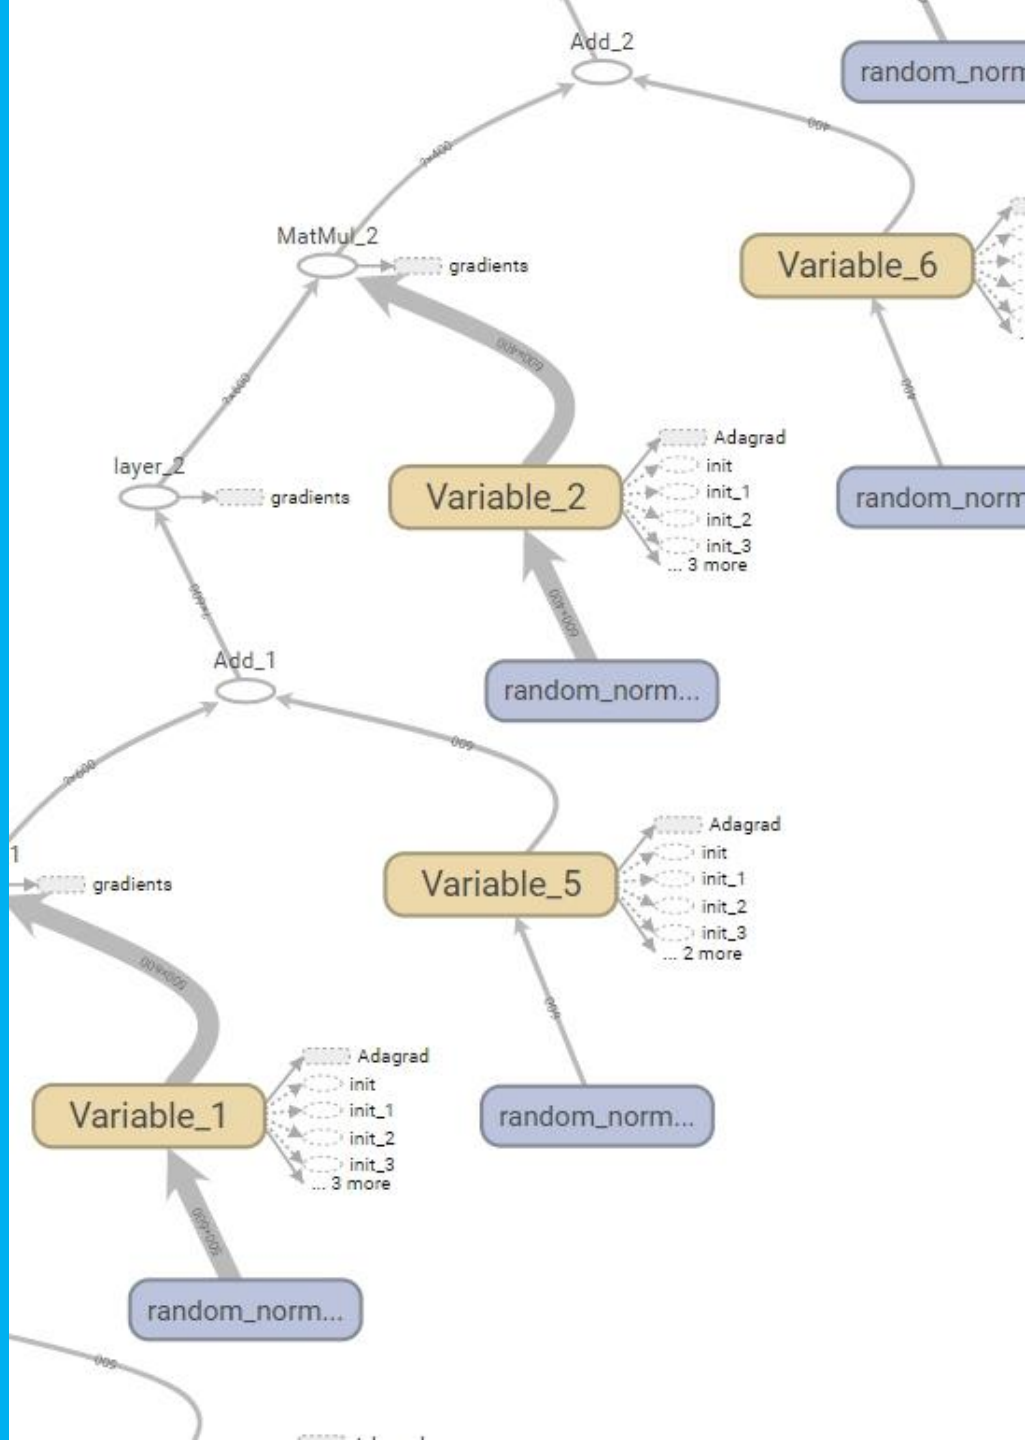
Huy Viet Le, Sven Mayer, Niels Henze



interactionlab.io

# Core Graph Data Structures

- `tf.Tensor`
  - Can be a `tf.variable`, `tf.constant`, `tf.placeholder`, ...
  - Example: `c = tf.constant([1,2,3])`
- `tf.Operation`
  - Node in a graph that take `tf.Tensor(s)` as input and return `tf.Tensor(s)` as output.
  - Example: `c = tf.matmul(a, b)`
- `tf.Graph`
  - Contain set of `tf.Operation` and `tf.Tensor` object for using them e.g. in sessions.



# Data Structures in TensorFlow

## How models are stored

- `NodeDef`: Operations and Fixed Values
  - A single *operation* in a model and the basic unit of computation in TensorFlow.
  - Operation type (e.g. Add, or Mul) and parameters needed to execute the operation.
  - `Const` are also stored in a `NodeDef`
- `GraphDef`
  - Contains a list of `NodeDefs` to define the computational graph.
  - Results in `.pbtxt` (text representation) or `.pb` (binary representation) files.
- `MetaGraphDef`
  - Contains a `GraphDef` and 'signatures' (e.g. checkpoint paths, and I/O node information).
  - Results in `.meta` files and required for re-training.

# Data Structures in TensorFlow

## How models are stored

- `Checkpoint`: Variable Values
  - Stores *variables* (whose values are held in RAM for time-critical operations).
  - Examples are weights and biases.
  - Results in `*.data` files.
- `SavedModel`: A bundle of all required data
  - Includes `MetaGraphDef`, checkpoint files, and further asset files (e.g. label names).

# Porting Models to Mobile and Embedded Devices

## 'Freezing' graphs

- Running a computational graph requires `NodeDefs` as well as `Variables`.
- `Variables` (e.g. weights) are stored in `Checkpoints` and not in `GraphDefs`.
- Inconvenient to have separate files for deployment in production.
- Freezing: Converting all `Variables` to `Const` to embed them in a single `GraphDef`.
  - Results in `'.pbtxt'` (text format) or `'pb'` (binary format).

# Using Models on an Android Device

## Overview of procedure

- Add the TensorFlow Inference Interface into the gradle build file:  
`'org.tensorflow:tensorflow-android:+'`
- Copy model file (\*.pb) to the project's assets folder.
- Initialize `TensorFlowInferenceInterface` with the model file.
- Specify input/output node names, and input data structure.
- Retrieve inference result and decode one-hot encoding (for classification).

# Live Code Demo

# **Optimizing Models for Mobile Devices**



# Graph Transform Tool

- Needs to be manually built with Bazel (<https://bazel.build/>).
- In case TensorFlow is already built with Bazel:
  - Build: `bazel build tensorflow/tools/graph_transforms:transform_graph`
  - Run: `bazel run tensorflow/tools/graph:transforms:transform_graph --`
- In case TensorFlow was installed with pip:
  - Download TensorFlow sources from github and unpack.
  - Build: `bazel build tensorflow/tools/graph_transforms:transform_graph`
  - SymLink to the executables will be created in the root folder.
  - Run: `bazel-bin/tensorflow/tools/graph_transforms/transform_graph --`

# Getting an Overview of the Graph

## Graph Summary & Benchmark

```
bazel-bin/tensorflow/tools/graph_transforms/summarize_graph  
  --in_graph=original_model.pb
```

# Optimizing the Graph for Mobile Devices

## Quantization of Neural Networks

- Weights and biases are stored as 32-bit floating points.
- Idea: Reducing floating point accuracy to eight-bit equivalent for constants.
- 70% compression in comparison to 8% using gzip [1].
  - Accuracy Loss: ~1% [1].

```
bazel-bin/tensorflow/tools/graph_transforms/transform_graph
  --in_graph=original_model.pb
  --out_graph=quantized_model.pb
  --inputs=input
  --outputs=output
  --transforms='quantize_weights'
```

[1] Source: <https://codelabs.developers.google.com/codelabs/tensorflow-for-poets-2/#4>

# Optimizing the Graph for Mobile Devices

## Removing unused nodes

- Remove `GraphDefs` contain `NodeDefs` that are needed for e.g. summary ops.
  - For inference only, these `NodeDefs` are not needed.

```
bazel-bin/tensorflow/tools/graph_transforms/transform_graph\  
  --in_graph=original_model.pb\  
  --out_graph=optimized_model.pb\  
  --inputs=input\  
  --outputs=output\  
  --transforms='strip_unused_nodes(type=float, shape="1,784")'
```

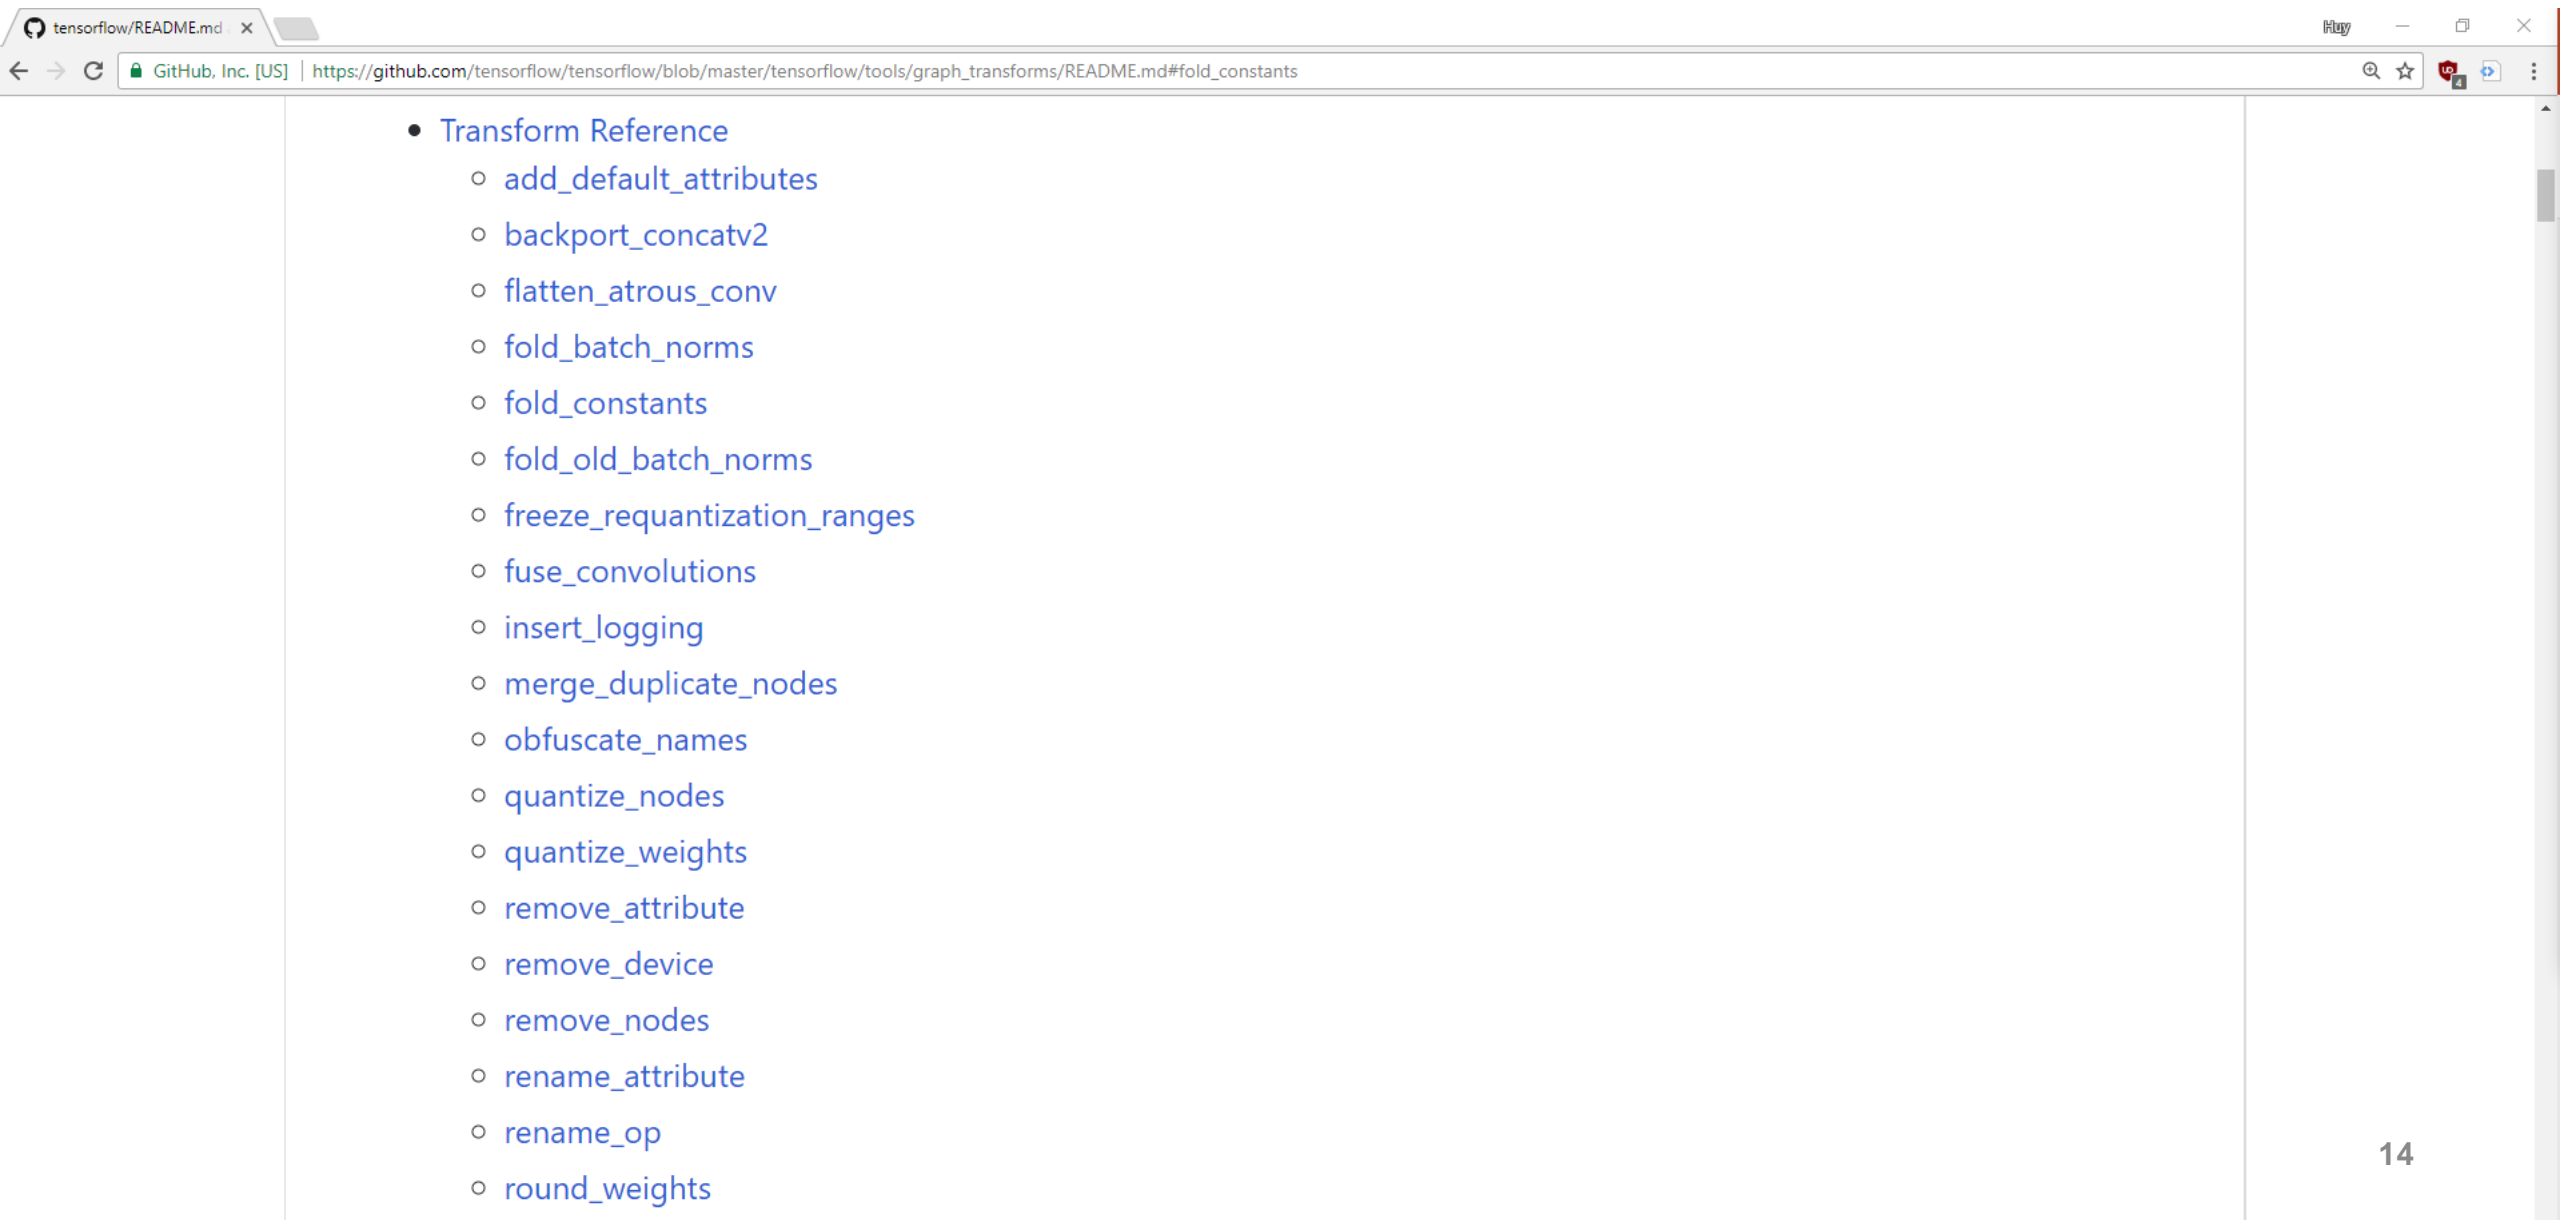
# Optimizing the Graph for Mobile Devices

## Fold constants

- Replaces sub-graphs that always evaluate to constant expressions with actual constants.

```
bazel-bin/tensorflow/tools/graph_transforms/transform_graph\  
  --in_graph=original_model.pb\  
  --out_graph=optimized_model.pb\  
  --inputs=input\  
  --outputs=output\  
  --transforms='fold_constants'
```

# Further Transformations



The image is a screenshot of a web browser window. The browser's address bar shows the URL: [https://github.com/tensorflow/tensorflow/blob/master/tensorflow/tools/graph\\_transforms/README.md#fold\\_constants](https://github.com/tensorflow/tensorflow/blob/master/tensorflow/tools/graph_transforms/README.md#fold_constants). The page content displays a list of transformations under the heading "Transform Reference".

- Transform Reference
  - [add\\_default\\_attributes](#)
  - [backport\\_concatv2](#)
  - [flatten\\_atrous\\_conv](#)
  - [fold\\_batch\\_norms](#)
  - [fold\\_constants](#)
  - [fold\\_old\\_batch\\_norms](#)
  - [freeze\\_requantization\\_ranges](#)
  - [fuse\\_convolutions](#)
  - [insert\\_logging](#)
  - [merge\\_duplicate\\_nodes](#)
  - [obfuscate\\_names](#)
  - [quantize\\_nodes](#)
  - [quantize\\_weights](#)
  - [remove\\_attribute](#)
  - [remove\\_device](#)
  - [remove\\_nodes](#)
  - [rename\\_attribute](#)
  - [rename\\_op](#)
  - [round\\_weights](#)